

STATS 3DA3

Homework Assignment 6

Shujia Zhang(400390028), Yunxin Li(400323756), Ling Dai(400220798)

2024-04-18

(Rubini and Eswaran 2015)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error, confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from patsy import dmatrices, dmatrix
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import accuracy_score
import statsmodels.api as sm
```

```
pip install ucimlrepo
```

Requirement already satisfied: ucimlrepo in /Users/dailing/Desktop/3da-assignment6/.conda/lib/

Note: you may need to restart the kernel to use updated packages.

```
from ucimlrepo import fetch_ucirepo

# fetch dataset
chronic_kidney_disease = fetch_ucirepo(id=336)

# data (as pandas dataframes)
X = chronic_kidney_disease.data.features
y = chronic_kidney_disease.data.targets

# metadata
```

```
print(chronic_kidney_disease.metadata)
```

```
# variable information
```

```
print(chronic_kidney_disease.variables)
```

```
{'uci_id': 336, 'name': 'Chronic Kidney Disease', 'repository_url': 'https://archive.ics.uci.edu/dataset/336/chronic_kidney_disease'}
```

	name	role	type	demographic	description \
0	age	Feature	Integer	Age	None
1	bp	Feature	Integer	None	blood pressure
2	sg	Feature	Categorical	None	specific gravity
3	al	Feature	Categorical	None	albumin
4	su	Feature	Categorical	None	sugar
5	rbc	Feature	Binary	None	red blood cells
6	pc	Feature	Binary	None	pus cell
7	pcc	Feature	Binary	None	pus cell clumps
8	ba	Feature	Binary	None	bacteria
9	bgr	Feature	Integer	None	blood glucose random
10	bu	Feature	Integer	None	blood urea
11	sc	Feature	Continuous	None	serum creatinine
12	sod	Feature	Integer	None	sodium
13	pot	Feature	Continuous	None	potassium
14	hemo	Feature	Continuous	None	hemoglobin
15	pcv	Feature	Integer	None	packed cell volume
16	wbcc	Feature	Integer	None	white blood cell count
17	rbcc	Feature	Continuous	None	red blood cell count
18	htn	Feature	Binary	None	hypertension
19	dm	Feature	Binary	None	diabetes mellitus
20	cad	Feature	Binary	None	coronary artery disease
21	appet	Feature	Binary	None	appetite
22	pe	Feature	Binary	None	pedal edema
23	ane	Feature	Binary	None	anemia
24	class	Target	Binary	None	ckd or not ckd

	units	missing_values
0	year	yes
1	mm/Hg	yes
2	None	yes
3	None	yes
4	None	yes
5	None	yes
6	None	yes
7	None	yes
8	None	yes
9	mgs/dl	yes
10	mgs/dl	yes
11	mgs/dl	yes
12	mEq/L	yes
13	mEq/L	yes
14	gms	yes
15	None	yes
16	cells/cmm	yes
17	millions/cmm	yes
18	None	yes
19	None	yes
20	None	yes
21	None	yes
22	None	yes
23	None	yes
24	None	no

X

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	hemo	pcv	w
0	48.0	80.0	1.020	1.0	0.0	NaN	normal	notpresent	notpresent	121.0	...	15.4	44.0	73
1	7.0	50.0	1.020	4.0	0.0	NaN	normal	notpresent	notpresent	NaN	...	11.3	38.0	60
2	62.0	80.0	1.010	2.0	3.0	normal	normal	notpresent	notpresent	423.0	...	9.6	31.0	73
3	48.0	70.0	1.005	4.0	0.0	normal	abnormal	present	notpresent	117.0	...	11.2	32.0	67
4	51.0	80.0	1.010	2.0	0.0	normal	normal	notpresent	notpresent	106.0	...	11.6	35.0	73
...
395	55.0	80.0	1.020	0.0	0.0	normal	normal	notpresent	notpresent	140.0	...	15.7	47.0	67
396	42.0	70.0	1.025	0.0	0.0	normal	normal	notpresent	notpresent	75.0	...	16.5	54.0	73
397	12.0	80.0	1.020	0.0	0.0	normal	normal	notpresent	notpresent	100.0	...	15.8	49.0	60
398	17.0	60.0	1.025	0.0	0.0	normal	normal	notpresent	notpresent	114.0	...	14.2	51.0	72
399	58.0	80.0	1.025	0.0	0.0	normal	normal	notpresent	notpresent	131.0	...	15.8	53.0	63

y

	class
0	ckd
1	ckd
2	ckd
3	ckd
4	ckd
...	...
395	notckd
396	notckd
397	notckd
398	notckd
399	notckd

X.dtypes

age float64

```
bp      float64
sg      float64
al      float64
su      float64
rbc     object
pc      object
pcc     object
ba      object
bgr     float64
bu      float64
sc      float64
sod     float64
pot     float64
hemo    float64
pcv     float64
wbcc    float64
rbcc    float64
htn     object
dm      object
cad     object
appet   object
pe      object
ane     object
dtype: object
```

```
y.dtypes
```

```
class    object
dtype: object
```

```
#to see if "rbc" is binary, since from the website information I found that this variable should be binary
X['rbc'].value_counts()
```

```
rbc
normal      201
abnormal    47
Name: count, dtype: int64
```

```
#to see if "pc" is binary, since from the website inormation I found that this variable should
X['pc'].value_counts()
```

```
pc
normal      259
abnormal    76
Name: count, dtype: int64
```

```
#to see if "pcc" is binary, since from the website inormation I found that this variable should
X['pcc'].value_counts()
```

```
pcc
notpresent  354
present     42
Name: count, dtype: int64
```

```
#to see if "ba" is binary, since from the website inormation I found that this variable should
X['ba'].value_counts()
```

```
ba
notpresent  374
present     22
Name: count, dtype: int64
```

```
#to see if "htn" is binary, since from the website inormation I found that this variable should
X['htn'].value_counts()
```

```
htn
no      251
yes     147
Name: count, dtype: int64
```

```
#to see if "dm" is binary, since from the website information I found that this variable should
X['dm'].value_counts()
```

```
dm
no      260
yes     137
\tno     1
Name: count, dtype: int64
```

```
#find that "dm" is not binary, but it should be binary.
X_cat = set(X["dm"])
X_cat
```

```
{'\tno', nan, 'no', 'yes'}
```

```
#make "dm" become binary
X['dm'] = X['dm'].replace('\tno', 'no')
X_cat = set(X["dm"])
X_cat
```

```
{nan, 'no', 'yes'}
```

```
#to double check if "dm" is binary now.
X['dm'].value_counts()
```

```
dm
no      261
yes     137
Name: count, dtype: int64
```



```
#to see if "cad" is binary, since from the website information I found that this variable should be binary
X['cad'].value_counts()
```

```
cad
no      364
yes      34
Name: count, dtype: int64
```

```
#to see if "appet" is binary, since from the website information I found that this variable should be binary
X['appet'].value_counts()
```

```
appet
good    317
poor     82
Name: count, dtype: int64
```

```
#to see if "pe" is binary, since from the website information I found that this variable should be binary
X['pe'].value_counts()
```

```
pe
no      323
yes      76
Name: count, dtype: int64
```

```
#to see if "ane" is binary, since from the website information I found that this variable should be binary
X['ane'].value_counts()
```

```
ane
no      339
yes      60
Name: count, dtype: int64
```

```
#to see if "class" is binary, since from the website information I found that this variable should be binary.
y['class'].value_counts()
```

```
class
ckd      248
notckd   150
ckd\t      2
Name: count, dtype: int64
```

```
#find "class" is not binary, but it should be binary.
y_cat = set(y["class"])
y_cat
```

```
{'ckd', 'ckd\t', 'notckd'}
```

```
# make it become binary
y['class'] = y['class'].replace('ckd\t', 'ckd')
y_cat = set(y["class"])
y_cat
```

```
{'ckd', 'notckd'}
```

```
#To double check if it is binary now.
y['class'].value_counts()
```

```
class
ckd      250
notckd   150
Name: count, dtype: int64
```

```
# number of missing values in X
X.isnull().sum()
```

```
age      9
bp       12
sg       47
al       46
su       49
rbc      152
pc       65
pcc      4
ba       4
bgr      44
bu       19
sc       17
sod      87
pot      88
hemo     52
pcv      71
wbcc     106
rbcc     131
htn      2
dm       2
cad      2
appet    1
pe       1
ane      1
dtype: int64
```

```
# number of missing values in y
y.isnull().sum()
```

```
class      0
dtype: int64
```

1. **Classification Problem Identification:** Define and describe a classification problem based on the dataset.

There are 400 observations and 25 variables in Chronic Kidney Disease dataset. There are 14 float64 type variables, and they are “age”, “bp”, “sg”, “al”, “su”, “bgr”, “bu”, “sc”, “sod”, “pot”, “hemo”, “pcv”, “wbcc”, “rbcc”. There are 11 object type variables, and they are “rbc”, “pc”, “pcc”, “ba”, “htn”, “dm”, “cad”, “appet”, “pe”, “ane”, “calss”. In these 25 variables, all variables in X are covariates (consists 14 floats and 10 object) and the variable in y is the response (the variable “class” which is an object variable). Moreover, the response is “class” contains only 2 category: “ckd” or “not ckd”, so the response is binary.

There are 9 missing values in “age”, 12 missing values in “bp”, 47 missing values in “sg”, 46 missing values in “al”, 49 missing values in “su”, 152 missing values in “rbc”, 65 missing values in “pc”, 4 missing values in “pcc”, 4 missing values in “ba”, 44 missing values in “bgr”, 19 missing values in “bu”, 17 missing values in “sc”, 87 missing values in “sod”, 88 missing values in “pot”, 52 missing values in “hemo”, 71 missing values in “pcv”, 106 missing values in “wbcc”, 131 missing values in “rbcc”, 2 missing values in “htn”, 2 missing values in “dm”, 2 missing values in “cad”, 1 missing value in “appet”, 1 missing value in “pe” and 1 missing value in “ane”.

0 missing value in “class”.

2. **Variable Transformation:** Implement any transformations chosen or justify the absence of such modifications.

```
X1 = X.copy()
columns_to_convert = ['sg', 'al', 'su']
for col in columns_to_convert:
    X1[col] = pd.Categorical(X1[col])

object_convert = ["rbc", "pc", "pcc", "ba", "htn", "dm", "cad", "appet", "pe", "ane"]
for col in object_convert:
```

```
X1[col] = X1[col].astype('category').cat.codes
X1.dtypes
```

```
age      float64
bp       float64
sg       category
al       category
su       category
rbc      int8
pc       int8
pcc      int8
ba       int8
bgr      float64
bu       float64
sc       float64
sod      float64
pot      float64
hemo     float64
pcv      float64
wbcc     float64
rbcc     float64
htn      int8
dm       int8
cad      int8
appet    int8
pe       int8
ane      int8
dtype: object
```

```
#cat = ['sg', 'al', 'su', 'rbc', 'pc', 'pcc', 'ba', 'htn', 'dm', 'cad', 'appet', 'pe', 'ane']
#X = pd.get_dummies(X,columns=cat)
#caty = ['class']
```

```
#y = pd.get_dummies(y,columns=caty)
```

```
X1
```

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	hemo	pcv	wbcc	rbcc	htn	dm
0	48.0	80.0	1.020	1.0	0.0	-1	1	0	0	121.0	...	15.4	44.0	7800.0	5.2	1	1
1	7.0	50.0	1.020	4.0	0.0	-1	1	0	0	NaN	...	11.3	38.0	6000.0	NaN	0	0
2	62.0	80.0	1.010	2.0	3.0	1	1	0	0	423.0	...	9.6	31.0	7500.0	NaN	0	1
3	48.0	70.0	1.005	4.0	0.0	1	0	1	0	117.0	...	11.2	32.0	6700.0	3.9	1	0
4	51.0	80.0	1.010	2.0	0.0	1	1	0	0	106.0	...	11.6	35.0	7300.0	4.6	0	0
...
395	55.0	80.0	1.020	0.0	0.0	1	1	0	0	140.0	...	15.7	47.0	6700.0	4.9	0	0
396	42.0	70.0	1.025	0.0	0.0	1	1	0	0	75.0	...	16.5	54.0	7800.0	6.2	0	0
397	12.0	80.0	1.020	0.0	0.0	1	1	0	0	100.0	...	15.8	49.0	6600.0	5.4	0	0
398	17.0	60.0	1.025	0.0	0.0	1	1	0	0	114.0	...	14.2	51.0	7200.0	5.9	0	0
399	58.0	80.0	1.025	0.0	0.0	1	1	0	0	131.0	...	15.8	53.0	6800.0	6.1	0	0

```
X1['sg'].value_counts()
```

```
sg
```

```
1.020    106
```

```
1.010     84
```

```
1.025     81
```

```
1.015     75
```

```
1.005      7
```

```
Name: count, dtype: int64
```

```
#why 3 categories?
```

```
X1['rbc'].value_counts()
```

```
rbc
```

```
1    201
```

```
-1    152
0      47
```

Name: count, dtype: int64

```
X1['pe'].value_counts()
```

```
pe
0    323
1     76
-1      1
```

Name: count, dtype: int64

```
float_col = X1.select_dtypes(include=['float64']).columns
object_col = X1.select_dtypes(include=['object']).columns
```

```
#way2
```

```
binary = {'rbc':{'normal':1,'abnormal':0},
          'pc':{'normal':1,'abnormal':0},
          'pcc':{'present':1,'notpresent':0},
          'ba':{'present':1,'notpresent':0},
          'htn':{'yes':1,'no':0},
          'dm':{'yes':1,'no':0},
          'cad':{'yes':1,'no':0},
          'appet':{'good':1,'poor':0},
          'pe':{'yes':1,'no':0},
          'ane':{'yes':1,'no':0}

          }
```

```
for i,j in binary.items():
    X1[i] = X1[i].replace(j)
```

```
X1
```

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	hemo	pcv	wbcc	rbcc	htn	dm	...
0	48.0	80.0	1.020	1.0	0.0	-1	1	0	0	121.0	...	15.4	44.0	7800.0	5.2	1	1	...
1	7.0	50.0	1.020	4.0	0.0	-1	1	0	0	NaN	...	11.3	38.0	6000.0	NaN	0	0	...
2	62.0	80.0	1.010	2.0	3.0	1	1	0	0	423.0	...	9.6	31.0	7500.0	NaN	0	1	...
3	48.0	70.0	1.005	4.0	0.0	1	0	1	0	117.0	...	11.2	32.0	6700.0	3.9	1	0	...
4	51.0	80.0	1.010	2.0	0.0	1	1	0	0	106.0	...	11.6	35.0	7300.0	4.6	0	0	...
...
395	55.0	80.0	1.020	0.0	0.0	1	1	0	0	140.0	...	15.7	47.0	6700.0	4.9	0	0	...
396	42.0	70.0	1.025	0.0	0.0	1	1	0	0	75.0	...	16.5	54.0	7800.0	6.2	0	0	...
397	12.0	80.0	1.020	0.0	0.0	1	1	0	0	100.0	...	15.8	49.0	6600.0	5.4	0	0	...
398	17.0	60.0	1.025	0.0	0.0	1	1	0	0	114.0	...	14.2	51.0	7200.0	5.9	0	0	...
399	58.0	80.0	1.025	0.0	0.0	1	1	0	0	131.0	...	15.8	53.0	6800.0	6.1	0	0	...

```
scale = StandardScaler()
X1[float_col] = scale.fit_transform(X1[float_col])
X1
```

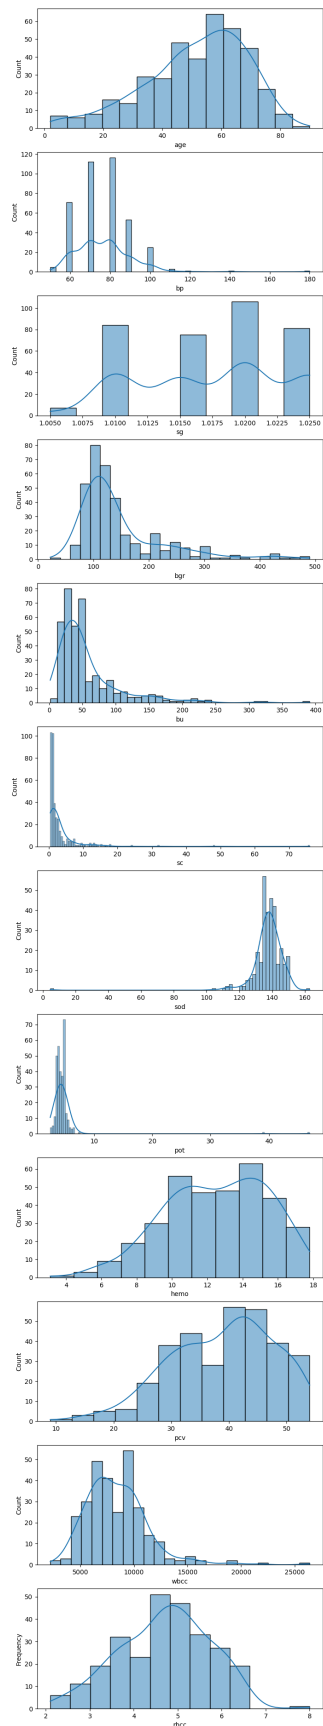
	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	hemo	pcv	wbcc	rbcc	htn	dm	...
0	-0.203139	0.258373	1.020	1.0	0.0	-1	1	0	0	-0.341498	...	0.988022	0.569881	-0.215188	5.2	1	1	...
1	-2.594124	-1.936857	1.020	4.0	0.0	-1	1	0	0	NaN	...	-0.421688	-0.098536	-0.815385	NaN	0	0	...
2	0.613295	0.258373	1.010	2.0	3.0	1	1	0	0	3.473064	...	-1.006202	-0.878356	-0.315385	NaN	0	1	...
3	-0.203139	-0.473370	1.005	4.0	0.0	1	0	1	0	-0.392022	...	-0.456071	-0.766953	-0.515385	3.9	1	0	...
4	-0.028189	0.258373	1.010	2.0	0.0	1	1	0	0	-0.530963	...	-0.318538	-0.432744	-0.315385	4.6	0	0	...
...
395	0.205078	0.258373	1.020	0.0	0.0	1	1	0	0	-0.101509	...	1.091172	0.904090	-0.515385	4.9	0	0	...
396	-0.553039	-0.473370	1.025	0.0	0.0	1	1	0	0	-0.922524	...	1.366237	1.683910	-0.215188	6.2	0	0	...
397	-2.302541	0.258373	1.020	0.0	0.0	1	1	0	0	-0.606749	...	1.125555	1.126896	-0.606749	5.4	0	0	...
398	-2.010957	-1.205114	1.025	0.0	0.0	1	1	0	0	-0.429915	...	0.575424	1.349701	-0.429915	5.9	0	0	...
399	0.380028	0.258373	1.025	0.0	0.0	1	1	0	0	-0.215188	...	1.125555	1.572507	-0.515385	6.1	0	0	...

3. **Dataset Overview:** Provide a detailed description of the dataset, covering variables, sum-

maries, observation counts, data types, and distributions (at least three statements).

```
def plot_dist(data, features):
    fig, axes = plt.subplots(nrows=len(features), ncols=1, figsize=(8, 4*len(features)))
    if len(features) == 1:
        axes = [axes]
    for ax, feature in zip(axes, features):
        sns.histplot(data[feature], kde=True, ax=ax)
    ax.set_xlabel(feature)
    ax.set_ylabel('Frequency')

numerical_cols = ["age", "bp", "sg", "bgr", "bu", "sc", "sod", "pot", "hemo", "pcv", "wbcc", "rbcc"]
plt.show()
plot_dist(X, numerical_cols)
```



```
X1.describe()
```

	age	bp	rbc	pc	pcc	ba	bgr	bu
count	3.910000e+02	3.880000e+02	400.00000	400.000000	400.000000	400.000000	3.560000e+02	3.910000e+02
mean	9.994847e-17	-2.380684e-16	0.12250	0.485000	0.095000	0.045000	-1.796316e-16	-3.414983e-01
std	1.001281e+00	1.001291e+00	0.93256	0.759089	0.325946	0.251262	1.001407e+00	1.001281e+00
min	-2.885708e+00	-1.936857e+00	-1.00000	-1.000000	-1.000000	-1.000000	-1.591967e+00	-1.591967e+00
25%	-5.530393e-01	-4.733701e-01	-1.00000	0.000000	0.000000	0.000000	-6.193803e-01	-6.193803e-01
50%	2.050779e-01	2.583733e-01	1.00000	1.000000	0.000000	0.000000	-3.414983e-01	-3.414983e-01
75%	7.590867e-01	2.583733e-01	1.00000	1.000000	0.000000	0.000000	1.890038e-01	1.890038e-01
max	2.246163e+00	7.575807e+00	1.00000	1.000000	1.000000	1.000000	4.319341e+00	6.193803e-01

```
X1.dtypes
```

```
age      float64
bp      float64
sg      category
al      category
su      category
rbc      int8
pc      int8
pcc      int8
ba      int8
bgr      float64
bu      float64
sc      float64
sod      float64
pot      float64
hemo     float64
pcv      float64
wbcc     float64
```

```

rbcc      float64
htn       int8
dm        int8
cad       int8
appet     int8
pe        int8
ane       int8
dtype: object

```

Since there are 25 variables and 400 observations in Chronic Kidney Disease dataset.

1). There are 14 float64 type variables, and they are “age”, “bp”, “sg”, “al”, “su”, “bgr”, “bu”, “sc”, “sod”, “pot”, “hemo”, “pcv”, “wbcc”, “rbcc”. There are 11 object type variables, and they are “rbc”, “pc”, “pcc”, “ba”, “htn”, “dm”, “cad”, “appet”, “pe”, “ane”, “calss”. In these 25 variables, all variables in X are covariates (consists 14 floats and 10 object) and the variable in y is the response (the variable “class” which is an object variable). Moreover, the response is “class” contains only 2 category: “ckd” or “not ckd”, so the response is binary.

2). The number of features is 24, and they are: “age”, “bp”, “sg”, “al”, “su”, “bgr”, “bu”, “sc”, “sod”, “pot”, “hemo”, “pcv”, “wbcc”, “rbcc”, “rbc”, “pc”, “pcc”, “ba”, “htn”, “dm”, “cad”, “appet”, “pe”, “ane”, “calss”.

3). The number of observations is 400. But we can see that the count of some variables is less than 400, for example: the count of age is 391 which is less than 400. This means there are missing values in some variables in this dataset.

4). The type of variables: There are 14 float64 type variables, and they are “age”, “bp”, “sg”, “al”, “su”, “bgr”, “bu”, “sc”, “sod”, “pot”, “hemo”, “pcv”, “wbcc”, “rbcc”. There are 11 object type variables, and they are “rbc”, “pc”, “pcc”, “ba”, “htn”, “dm”, “cad”, “appet”, “pe”, “ane”, “calss”. In these 25 variables, all variables in X are covariates (consists 14 floats and 10 object) and the variable in y is the response (the variable “class” which is an object variable). Moreover, the response is “class” contains only 2 category: “ckd” or “not ckd”, so the response is binary.


5). Missing values: There are 9 missing values in “age”, 12 missing values in “bp”, 47 missing values in “sg”, 46 missing values in “al”, 49 missing values in “su”, 152 missing values in “rbc”, 65 missing

values in “pc”, 4 missing values in “pcc”, 4 missing values in “ba”, 44 missing values in “bgr”, 19 missing values in “bu”, 17 missing values in “sc”, 87 missing values in “sod”, 88 missing values in “pot”, 52 missing values in “hemo”, 71 missing values in “pcv”, 106 missing values in “wbcc”, 131 missing values in “rbcc”, 2 missing values in “htn”, 2 missing values in “dm”, 2 missing values in “cad”, 1 missing value in “appet”, 1 missing value in “pe” and 1 missing value in “ane”.

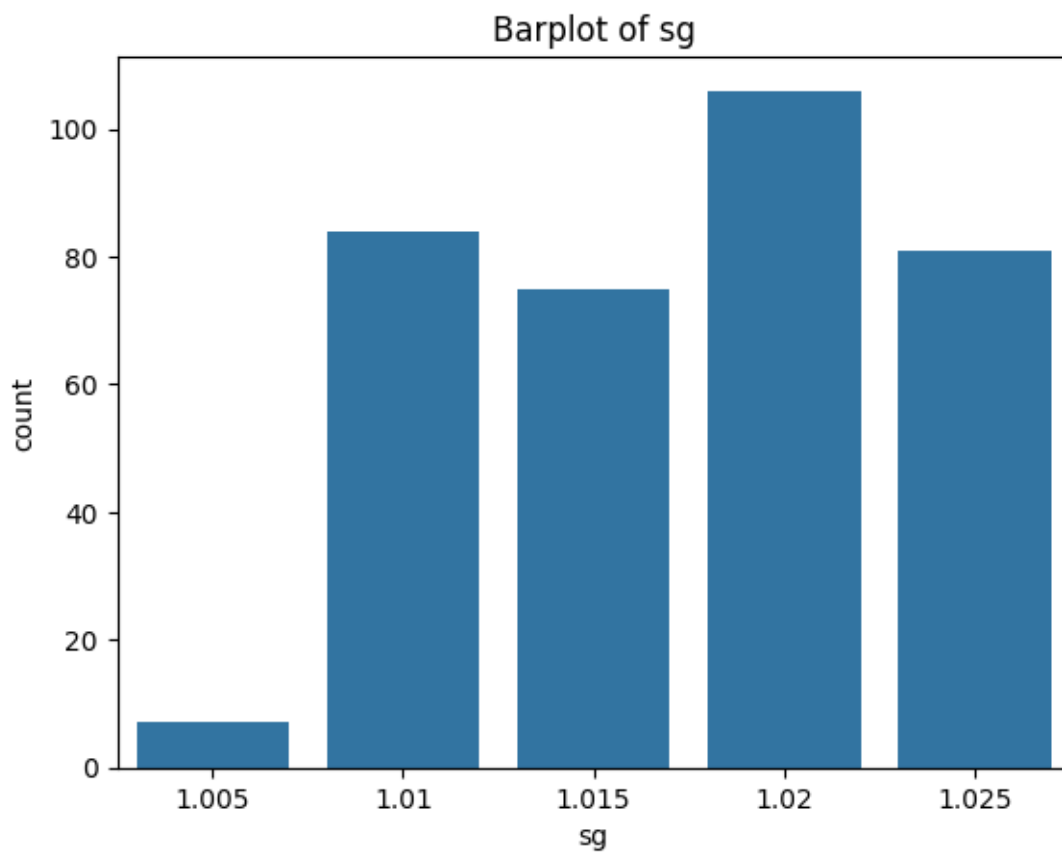
0 missing value in “class”.

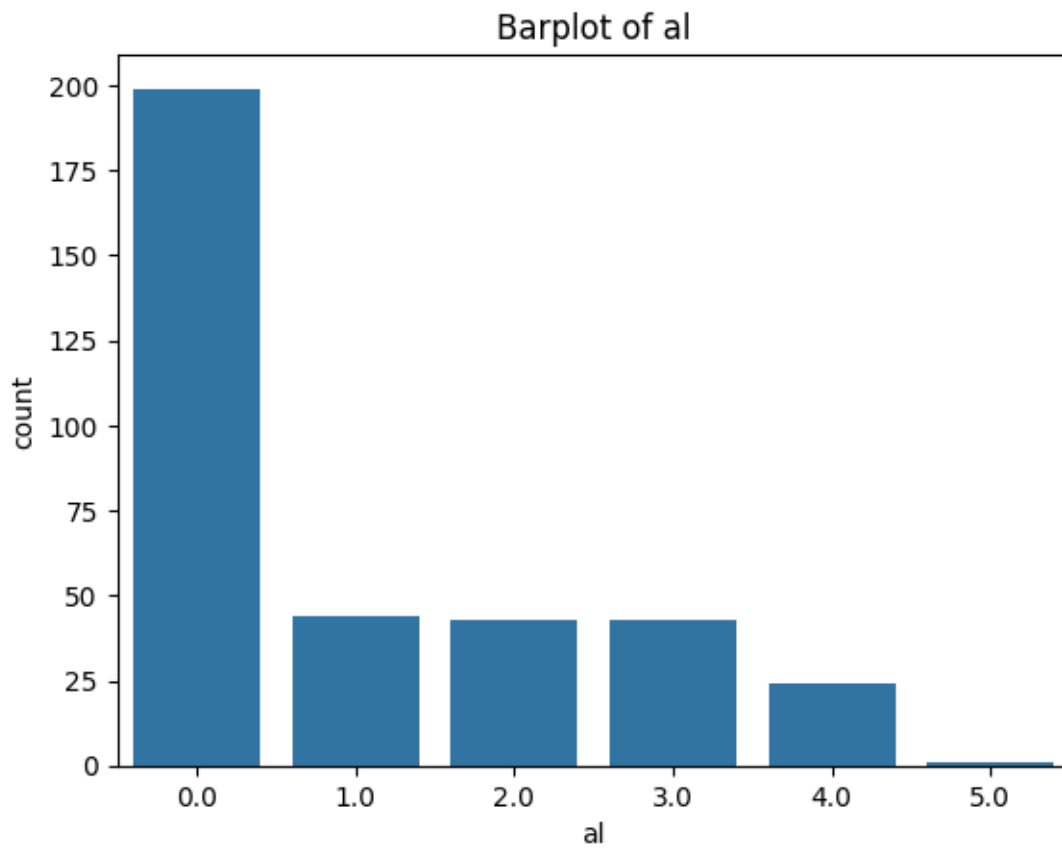
4. **Association Between Variables:** Analyze variable relationships and their implications for feature selection or extraction (at least three statements).


```
# prof says numerical variables shuold use heatmap, categorical variables should use barplot
#
#
#
#
#numerical variables:
float64=X1.select_dtypes(include=['float64'])
plt.figure(figsize=(15, 10))
sns.heatmap(float64.corr(),
            annot=True)
plt.title('Covariance Matrix Heatmap for numerical variables')
plt.show()
```





```
#categorical variables
categorical=X1.select_dtypes(include=['category','int8'])
for columnn in categorical.columns:
    plt.figure()
    sns.countplot(x=columnn, data=categorical)
    plt.title(f'Barplot of {columnn}')
    plt.show()
```




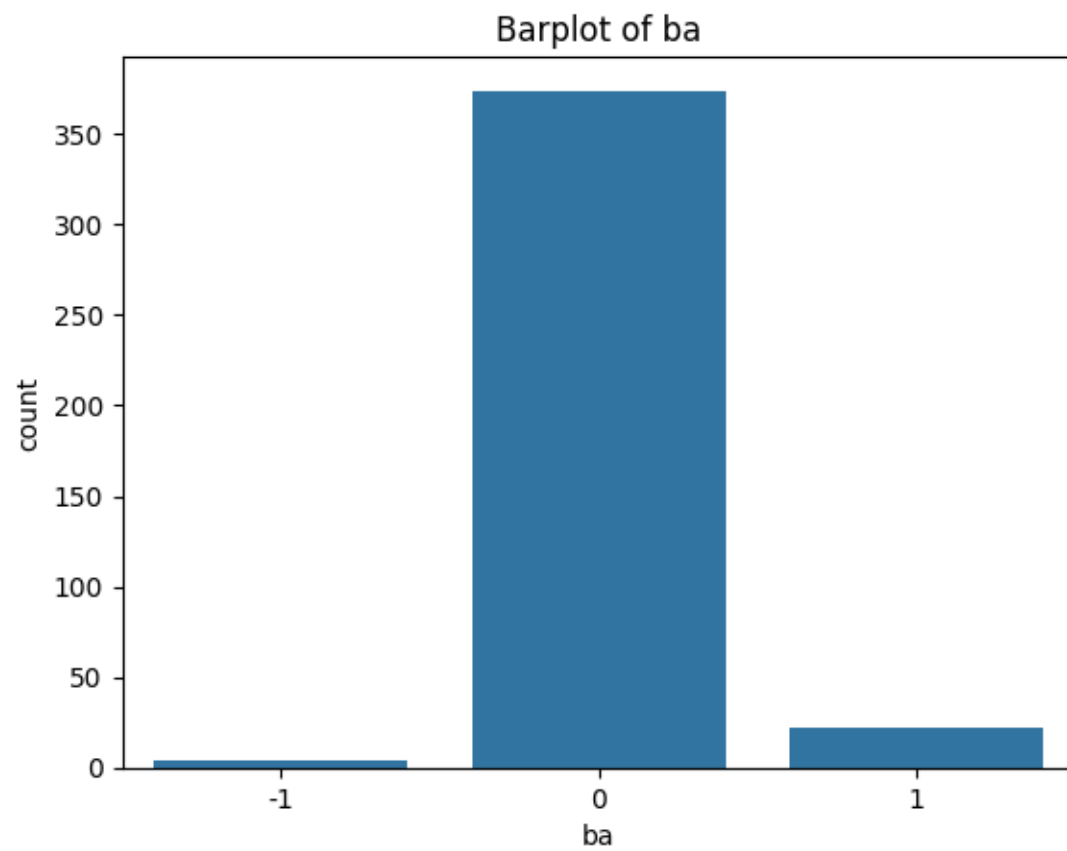





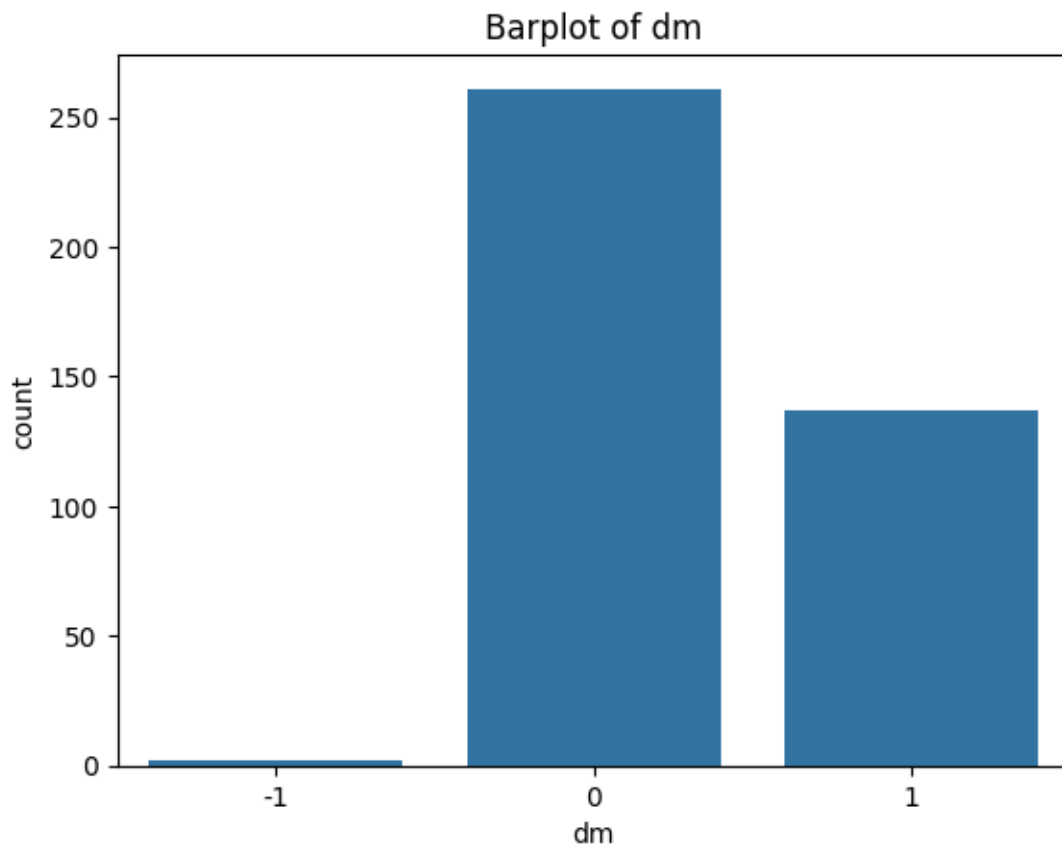


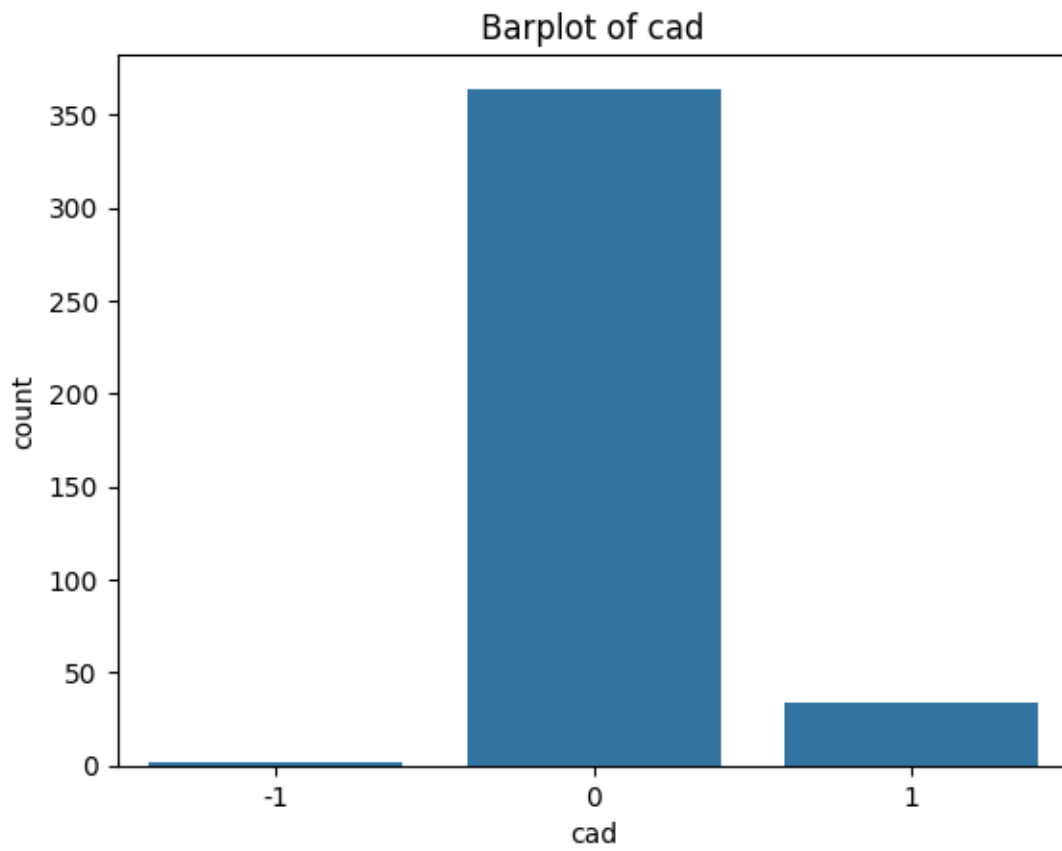





















We use heatmap for numerical variables and barplot for categorical variables.

For heatmap:

- 1). This is a symmetric matrix and the elements on the diagonal are the variance. The diagonal of the correlation matrix is 1.
- 2). The elements not on the diagonal are the covariance.
- 3). The smallest number in the matrix is -0.69 which means "sod" and "sc" are negatively correlated.
- 4). The largest number in the matrix is 1 (the diagonal).
- 5). The largest number other than 1 is 0.9 which means "hemo" and "pcv" are positively correlated.

For barplots:

- 1). In those binary variables, the category "-1" could suggest missing values.

2). For example, in “pe” (pedal edema), the most common category is “0”, which could imply the absence of pedal edema, and the second most common category is “1”, which could imply the presence of pedal edema, and the least category “-1” are missing values.

3). “htn” (hypertension) and “dm” (diabetes mellitus) are not good since the diistribution is not extreme.

5. Missing Value Analysis and Handling: Implement your strategy for identifying and addressing missing values in the dataset, or provide reasons for not addressing them.

```
X.isnull().sum()
```

age	9
bp	12
sg	47
al	46
su	49
rbc	152
pc	65
pcc	4
ba	4
bgr	44
bu	19
sc	17
sod	87
pot	88
hemo	52
pcv	71
wbcc	106
rbcc	131
htn	2
dm	2
cad	2

```
appet      1
pe         1
ane        1
dtype: int64
```

```
y.isnull().sum()
```

```
class      0
dtype: int64
```

```
#way2
float_na=X1[float_col].isna()
float_nasum=float_na.sum()
print(float_nasum)

cat_na=X1[columns_to_convert].isna()
cat_nasum=cat_na.sum()
print(cat_nasum)

binary_na = (X1[object_convert] == -1)
binary_nasum=binary_na.sum()
print(binary_nasum)
```

```
age        9
bp         12
bgr        44
bu         19
sc         17
sod        87
pot        88
hemo       52
pcv        71
```

```

wbcc    106
rbcc    131
dtype: int64
sg      47
al      46
su      49
dtype: int64
rbc     152
pc       65
pcc      4
ba       4
htn      2
dm       2
cad      2
appet    1
pe       1
ane      1
dtype: int64

```

Missing values: There are 9 missing values in “age”, 12 missing values in “bp”, 47 missing values in “sg”, 46 missing values in “al”, 49 missing values in “su”, 152 missing values in “rbc”, 65 missing values in “pc”, 4 missing values in “pcc”, 4 missing values in “ba”, 44 missing values in “bgr”, 19 missing values in “bu”, 17 missing values in “sc”, 87 missing values in “sod”, 88 missing values in “pot”, 52 missing values in “hemo”, 71 missing values in “pcv”, 106 missing values in “wbcc”, 131 missing values in “rbcc”, 2 missing values in “htn”, 2 missing values in “dm”, 2 missing values in “cad”, 1 missing value in “appet”, 1 missing value in “pe” and 1 missing value in “ane”.

0 missing value in “class”.

```

#the stragegy for addressing missing values
X1_drop= X1.dropna()
X1_drop

```

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	...	hemo	pcv	wbc
3	-0.203139	-0.473370	1.005	4.0	0.0	1	0	1	0	-0.392022	...	-0.456071	-0.766953	-0.5
5	0.496661	0.990117	1.015	3.0	0.0	-1	-1	0	0	-0.935155	...	-0.112239	0.012867	-0.2
9	0.088445	0.990117	1.020	2.0	0.0	0	0	1	0	-0.985679	...	-1.040585	-1.101161	1.23
11	0.671612	-0.473370	1.010	3.0	0.0	0	0	1	0	2.929931	...	-0.593604	-0.766953	-1.3
12	0.963195	-0.473370	1.015	3.0	1.0	-1	1	1	0	0.757399	...	-0.971819	-1.212564	1.29
...
395	0.205078	0.258373	1.020	0.0	0.0	1	1	0	0	-0.101509	...	1.091172	0.904090	-0.5
396	-0.553039	-0.473370	1.025	0.0	0.0	1	1	0	0	-0.922524	...	1.366237	1.683910	-0.2
397	-2.302541	0.258373	1.020	0.0	0.0	1	1	0	0	-0.606749	...	1.125555	1.126896	-0.6
398	-2.010957	-1.205114	1.025	0.0	0.0	1	1	0	0	-0.429915	...	0.575424	1.349701	-0.4
399	0.380028	0.258373	1.025	0.0	0.0	1	1	0	0	-0.215188	...	1.125555	1.572507	-0.5

dropped the columns if the missing value of this variable is greater than 44.

```
missing_values_count = X1.isna().sum()
columns_to_drop = missing_values_count[missing_values_count > 50].index
columns_to_drop
```

```
Index(['sod', 'pot', 'hemo', 'pcv', 'wbcc', 'rbcc'], dtype='object')
```

```
X1_dropped = X1.drop(columns=columns_to_drop)
X1_dropped
```

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	bu	sc	htn	dr
0	-0.203139	0.258373	1.020	1.0	0.0	-1	1	0	0	-0.341498	-0.424804	-0.326574	1	1
1	-2.594124	-1.936857	1.020	4.0	0.0	-1	1	0	0	NaN	-0.781687	-0.396338	0	0
2	0.613295	0.258373	1.010	2.0	3.0	1	1	0	0	3.473064	-0.087748	-0.221928	0	1
3	-0.203139	-0.473370	1.005	4.0	0.0	1	0	1	0	-0.392022	-0.028268	0.126891	1	0
4	-0.028189	0.258373	1.010	2.0	0.0	1	1	0	0	-0.530963	-0.623073	-0.291692	0	0
...

	age	bp	sg	al	su	rbc	pc	pcc	ba	bgr	bu	sc	htn	dr
395	0.205078	0.258373	1.020	0.0	0.0	1	1	0	0	-0.101509	-0.167055	-0.448661	0	0
396	-0.553039	-0.473370	1.025	0.0	0.0	1	1	0	0	-0.922524	-0.523939	-0.326574	0	0
397	-2.302541	0.258373	1.020	0.0	0.0	1	1	0	0	-0.606749	-0.623073	-0.431220	0	0
398	-2.010957	-1.205114	1.025	0.0	0.0	1	1	0	0	-0.429915	-0.147229	-0.361456	0	0
399	0.380028	0.258373	1.025	0.0	0.0	1	1	0	0	-0.215188	-0.781687	-0.344015	0	0

```
numerical_cols = X1_dropped.select_dtypes(include=['number']).columns
column_means = X1_dropped[numerical_cols].mean()

X1_filled = X1_dropped.copy()
X1_filled[numerical_cols] = X1_filled[numerical_cols].fillna(column_means)
```

Replacing the missing value by the mean of the variables for numerical variables.

```
from sklearn.impute import SimpleImputer
```

```
obj_col = X1_filled.select_dtypes(include=['category']).columns.tolist()
cat_imputer = SimpleImputer(strategy='most_frequent',missing_values=np.nan)
X1_filled.loc[:, obj_col] = cat_imputer.fit_transform(X1_filled.loc[:, obj_col])
```

Replacing the missing value by the most frequency input of the variables for categorical variables.

```
X1_filled.isna().sum()
```

```
age      0
bp       0
sg       0
al       0
su       0
rbc      0
pc       0
```



```
pcc      0
ba       0
bgr      0
bu       0
sc       0
htn      0
dm       0
cad      0
appet    0
pe       0
ane      0
dtype: int64
```

The strategy we used for addressing the missing value is: 1. Replacing the missing value of the numerical variables by its variable means. 2. Using the most frequent input to replaced the missing value for categorical variables.



The reason we delete the variables which have missing value size greater than 50 is because that large missing value may indicate problems in data qualities. If we using the above strategy to addressing the missing value, the value we replaced may cause 100% accuracy when classified the model. Thus, the most value this variable has missed, the most inaccuracy that the variable will cause when we classified it. Also when the missing vlues is too large, lets say in “rbc”, there are 152 missing vlues out of 400, then even though we train them and use althorithm to work on them, we will learn nothing from it.



In some algorithms, the missing values can be seen as a seperate category. For instance, in decision tree, KNN, etc. Moreover, when using the KNN classifier, the code can not train on “NA” response values.



6. **Outlier Analysis:** Implement your approach for identifying and managing outliers, or provide reasons for not addressing them.



```
for col in float_col:
    plt.figure(figsize=(6, 4))
```



```
sns.boxplot(x=X1.select_dtypes(include=['float64'])[col])  
plt.xlabel(col)  
plt.show()
```














7. **Sub-group Analysis:** Explore potential sub-groups within the data, employing appropriate data science methods to find the sub-groups of patients and visualize the sub-groups. The sub-group analysis must not include the labels (for CKD patients and healthy controls).

```
#prof says use K-means
from sklearn.preprocessing import scale
from sklearn.decomposition import PCA, TruncatedSVD
from sklearn.cluster import KMeans
from scipy.cluster import hierarchy
from sklearn.cluster import AgglomerativeClustering
from sklearn.metrics import silhouette_samples, silhouette_score
from sklearn.metrics.cluster import rand_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import adjusted_rand_score
from sklearn.decomposition import PCA, TruncatedSVD, FactorAnalysis
import matplotlib.pyplot as plt
from sklearn.metrics import silhouette_score, silhouette_samples
from matplotlib import cm
```

```
y['class'] = y["class"].astype('category').cat.codes
y
```

	class
0	0
1	0
2	0
3	0
4	0
...	...
395	1
396	1
397	1
398	1
399	1

[illegible]


```

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1], dtype=int32)

```

```

# imputer = KNNImputer(n_neighbors=5)
# X_imputed = pd.DataFrame(imputer.fit_transform(X1), columns=X1.columns)
# X_imputed.head()
# X_imputed.isna().sum()

```

```

range_n_clusters = [2, 3, 4, 5, 6]
for n_clusters in range_n_clusters:
    km = KMeans(n_clusters=n_clusters, n_init=20, random_state=0)
    labels = km.fit_predict(X1_drop)

    silhouette_avg = silhouette_score(X1_drop, labels)
    sample_silhouette_values = silhouette_samples(X1_drop, labels)
    fig, ax1 = plt.subplots(1, 1)
    fig.set_size_inches(18, 7)
    ax1.set_xlim([-0.3, 1])

    y_lower = 10

    for i in range(n_clusters):
        ith_cluster_silhouette_values = sample_silhouette_values[labels == i]
        ith_cluster_silhouette_values.sort()

```

```

size_cluster_i = ith_cluster_silhouette_values.shape[0]
y_upper = y_lower + size_cluster_i

# Use the colormap for coloring
color = cm.nipy_spectral(float(i) / n_clusters)
ax1.fill_betweenx(
    y=np.arange(y_lower, y_upper),
    x1=0,
    x2=ith_cluster_silhouette_values,
    facecolor=color,
    edgecolor=color,
    alpha=0.7,
)

ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))




y_lower = y_upper + 10

ax1.set_title("The silhouette plot for various clusters")
ax1.set_xlabel("The silhouette coefficient values")
ax1.set_ylabel("Cluster label")

ax1.axvline(x=silhouette_avg, color="red", linestyle="--")
plt.title(
    "Silhouette analysis for KMeans clustering on sample data with n_clusters = %d"
    % n_clusters,
    fontsize=14,
    fontweight="bold",
)

plt.show()

```




```

0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1], dtype=int32)

```

```
pca_X = PCA()
```

```
X = pd.DataFrame(X1_filled, X1_filled.index, X1_filled.columns)
```

```

pca_loading = pd.DataFrame(pca_X.fit(X).components_.T, index=X1_filled.columns, columns=['pc1'
        'pc11', 'pc12', 'pc13', 'pc14', 'pc15', 'pc16', 'pc17', 'pc18'
        # , 'pc19', 'pc20', 'pc21', 'pc22', 'pc23', 'pc24']
        ])

```

```
pca_loading
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
age	0.227990	0.230990	0.295718	-0.031051	-0.575582	-0.656010	0.133337	0.002424	0.021785
bp	0.218657	0.195864	0.182461	-0.453550	0.717207	-0.375424	-0.036042	0.093023	0.039444
sg	-0.001662	0.000410	0.000035	0.000870	0.000296	-0.000717	-0.001237	0.000471	0.001026
al	0.553457	-0.487707	-0.482395	-0.347007	-0.140393	0.010337	0.180866	0.053118	0.087756
su	0.357932	-0.308756	0.426065	0.400909	0.206522	0.046562	0.154558	-0.475340	0.349520
rbc	-0.160721	-0.314032	-0.325148	0.422191	0.147226	-0.553442	-0.342586	0.238952	0.237721
pc	-0.174043	-0.281904	-0.051507	0.116840	0.165232	-0.277618	0.505481	-0.186919	-0.676205
pcc	0.071946	-0.044355	-0.032420	-0.030809	-0.042205	-0.030729	-0.036475	0.017117	0.033291
ba	0.043953	-0.029370	-0.034667	-0.031891	0.001018	-0.018845	-0.006731	-0.005783	0.039313

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
bgr	0.337320	-0.236555	0.377282	0.237826	0.014970	0.158208	-0.269555	0.548677	-0.422632
bu	0.348920	0.344499	-0.344170	0.250931	0.056552	-0.049455	-0.410254	-0.465145	-0.346444
sc	0.266765	0.457515	-0.264006	0.424086	0.151383	0.091202	0.523546	0.375210	0.133144
htn	0.187133	0.065908	0.030741	-0.085847	-0.078187	-0.009468	-0.061016	0.053573	-0.091387
dm	0.180008	0.035135	0.115632	-0.019354	-0.062214	0.046096	-0.091541	0.012088	-0.044199
cad	0.061433	0.013987	0.014001	0.013954	-0.026179	-0.008642	0.008890	-0.007925	-0.000104
appet	0.094301	0.028995	-0.029747	-0.083728	-0.020176	0.013642	-0.047902	0.015283	-0.123596
pe	0.099282	0.005992	-0.074849	-0.057811	-0.050081	0.052186	-0.058089	-0.086412	0.013031
ane	0.076863	0.059741	-0.053879	-0.025595	0.035715	0.027753	-0.130954	-0.035215	-0.097805

```
pc_scores = pd.DataFrame(pca_X.fit_transform(X), columns=['pc1', 'pc2', 'pc3', 'pc4', 'pc5', 'pc6', 'pc7', 'pc8', 'pc9', 'pc10', 'pc11', 'pc12', 'pc13', 'pc14', 'pc15', 'pc16', 'pc17', 'pc18'], index=X.index)
pc_scores
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
0	-0.159842	0.123121	0.340298	-1.076715	-0.028920	0.446685	0.604284	-0.200452	-0.546039
1	0.080324	-2.658949	-2.090773	-0.980571	-0.533556	2.899035	1.077587	0.010434	-0.315841
2	2.748405	-2.445865	1.908401	1.610708	0.460511	-0.398526	-0.557039	0.808465	-0.761516
3	1.597823	-1.458225	-2.420973	-0.957120	-0.769911	-0.046338	-0.110910	0.412311	0.653469
4	-0.378850	-1.060810	-0.899482	-0.559600	0.156071	-0.818241	0.426550	0.266685	0.309952
...
395	-1.170494	-0.047818	0.180807	0.277176	0.311051	-0.960873	-0.289116	0.125640	-0.220852
396	-1.872246	-0.239130	-0.396054	0.399566	0.208608	-0.289936	0.067806	-0.182926	0.220656
397	-2.067094	-0.656654	-0.599016	0.127848	1.723678	0.628361	-0.291070	0.061005	0.098355
398	-2.076334	-0.721928	-0.895292	0.973608	0.546345	0.997310	-0.405810	-0.172560	-0.183314
399	-1.355503	-0.144376	0.373565	0.134861	0.189736	-1.053690	0.071790	0.388850	0.057877

```
var=pc_scores.var()
var
```


```
pc1      3.011006
pc2      1.418726
pc3      1.376123
pc4      0.871790
pc5      0.843322
pc6      0.780234
pc7      0.413204
pc8      0.356221
pc9      0.333497
pc10     0.193763
pc11     0.135449
pc12     0.107061
pc13     0.098757
pc14     0.082265
pc15     0.071608
pc16     0.058557
pc17     0.044481
pc18     0.000018
dtype: float64
```

```
plt.figure(figsize=(7,5))

plt.plot([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18]
, pca_X.explained_variance_ratio_, '-o', label='Individual component')
plt.plot([1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18]
, np.cumsum(pca_X.explained_variance_ratio_), '-s', label='Cumulative')

plt.ylabel('Proportion of Variance Explained')
plt.xlabel('Principal Component')
plt.xlim(0.75,4.25)
plt.ylim(0,1.05)
plt.xticks([1,2,3,4,5,6])
```

```
plt.legend(loc=2);
```



```
pca_loading
```

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
age	0.227990	0.230990	0.295718	-0.031051	-0.575582	-0.656010	0.133337	0.002424	0.021785
bp	0.218657	0.195864	0.182461	-0.453550	0.717207	-0.375424	-0.036042	0.093023	0.039444
sg	-0.001662	0.000410	0.000035	0.000870	0.000296	-0.000717	-0.001237	0.000471	0.001026
al	0.553457	-0.487707	-0.482395	-0.347007	-0.140393	0.010337	0.180866	0.053118	0.087756
su	0.357932	-0.308756	0.426065	0.400909	0.206522	0.046562	0.154558	-0.475340	0.349520
rbc	-0.160721	-0.314032	-0.325148	0.422191	0.147226	-0.553442	-0.342586	0.238952	0.237721
pc	-0.174043	-0.281904	-0.051507	0.116840	0.165232	-0.277618	0.505481	-0.186919	-0.676205
pcc	0.071946	-0.044355	-0.032420	-0.030809	-0.042205	-0.030729	-0.036475	0.017117	0.033291
ba	0.043953	-0.029370	-0.034667	-0.031891	0.001018	-0.018845	-0.006731	-0.005783	0.039313
bgr	0.337320	-0.236555	0.377282	0.237826	0.014970	0.158208	-0.269555	0.548677	-0.422632

	pc1	pc2	pc3	pc4	pc5	pc6	pc7	pc8	pc9
bu	0.348920	0.344499	-0.344170	0.250931	0.056552	-0.049455	-0.410254	-0.465145	-0.346444
sc	0.266765	0.457515	-0.264006	0.424086	0.151383	0.091202	0.523546	0.375210	0.133144
htn	0.187133	0.065908	0.030741	-0.085847	-0.078187	-0.009468	-0.061016	0.053573	-0.091387
dm	0.180008	0.035135	0.115632	-0.019354	-0.062214	0.046096	-0.091541	0.012088	-0.044199
cad	0.061433	0.013987	0.014001	0.013954	-0.026179	-0.008642	0.008890	-0.007925	-0.000104
appet	0.094301	0.028995	-0.029747	-0.083728	-0.020176	0.013642	-0.047902	0.015283	-0.123596
pe	0.099282	0.005992	-0.074849	-0.057811	-0.050081	0.052186	-0.058089	-0.086412	0.013031
ane	0.076863	0.059741	-0.053879	-0.025595	0.035715	0.027753	-0.130954	-0.035215	-0.097805

```
PCA_var=sum(pca_X.explained_variance_ratio_[0:2])
PCA_var
```

0.4344542793757311

```
plt.figure(figsize=(8, 6))


principal_components = pca_X.fit_transform(X)
kmeans = KMeans(n_clusters=2, n_init=20, random_state=0)
kmeans.fit(principal_components)
cluster_labels = kmeans.labels_

# Plot data points
plt.scatter(principal_components[:, 0], principal_components[:, 1], c=cluster_labels, cmap='viridis')

# Plot cluster centers
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], marker='x', c='red', s=100)

plt.title('K-means Clustering Visualization (k=2)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
```

```
plt.legend()
plt.grid(True)
plt.show()
```



```
y_1d = np.ravel(y)

adjusted_Rand_index = adjusted_rand_score(y_1d, cluster_labels)

adjusted_Rand_index
```

0.2663417112554181

8. **Data Splitting:** Segregate 30% of the data for testing, using a random seed of 1. Use the remaining 70% for training and model selection.

```
X_train, X_test, y_train, y_test = train_test_split(
    X1_filled, y, test_size=0.3, random_state=0, stratify=y) #or test_size=0.7
```

9. **Classifier Choices:** Identify the two classifiers you have chosen and justify your selections.

Algorithm 1: Decision tree

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

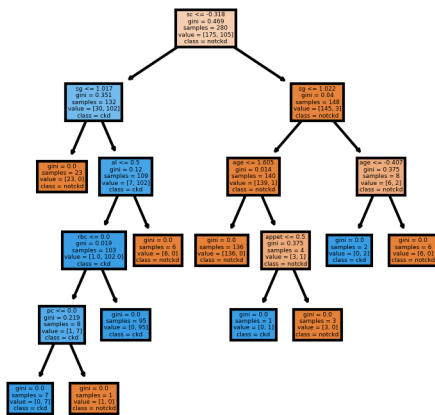
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error, confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
```

```
cs_dt = DecisionTreeClassifier(
    max_depth =30,
    random_state=1
)
```

```
cs_dt.fit(X_train, y_train)
```

```
DecisionTreeClassifier(max_depth=30, random_state=1)
```

```
fig, axes = plt.subplots(
    nrows = 1,ncols = 1,figsize = (3,3), dpi=300
)
plot_tree(
    cs_dt,
    max_depth= 30,
    feature_names = X_train.columns.tolist(),
```

```
pred = cs_dt.predict(X_test)
pred[:5]
```

```
array([0, 0, 0, 0, 0], dtype=int8)
```

```
cm = pd.DataFrame(confusion_matrix(y_test, pred), index=['No', 'Yes'], columns=['No', 'Yes'])
cm.index.name = 'True'
cm.columns.name = 'Predicted'
cm
```

	Predicted	
	No	Yes
True	No	Yes
	72	3
Yes	1	44

```
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.99	0.96	0.97	75
1	0.94	0.98	0.96	45

accuracy			0.97	120
macro avg	0.96	0.97	0.96	120
weighted avg	0.97	0.97	0.97	120

```
path = cs_dt.cost_complexity_pruning_path(
    X_train,
    y_train
)
ccp_alphas, impurities = path.ccp_alphas, path.impurities
```


```
clfs = []
for ccp_alpha in ccp_alphas:
    clf = DecisionTreeClassifier(
        random_state=0,
        ccp_alpha=ccp_alpha
    )
    clf.fit(X_train, y_train)
    clfs.append(clf)
```

```
depth = [clf.tree_.max_depth for clf in clfs]
depth
```

```
[5, 4, 3, 3, 2, 1, 0]
```

```
test_score = [clf.score(X_test, y_test) for clf in clfs]
```

```
plt.plot(depth, test_score)
plt.xlabel('Depth')
plt.ylabel('Accuracy')
plt.show()
```



```
cs_dt_new = DecisionTreeClassifier(  
    max_depth =3,  
    random_state=1  
)
```

```
cs_dt_new.fit(X_train, y_train)
```

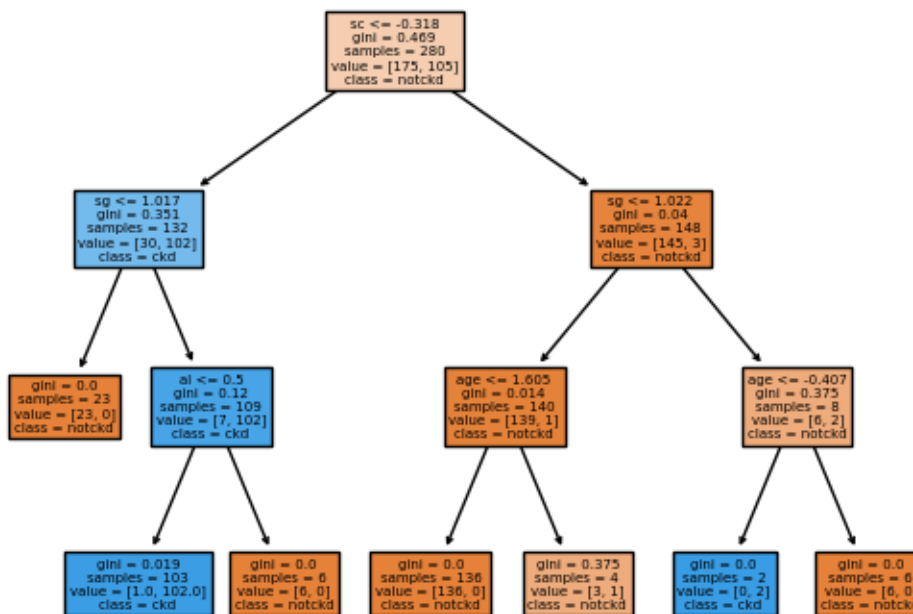
```
DecisionTreeClassifier(max_depth=3, random_state=1)
```

```
plot_tree(  
    cs_dt_new,  
    max_depth= 3,  
    feature_names = X_train.columns.tolist(),  
    class_names=['notckd', 'ckd'],  
    filled = True  
)
```

```

[Text(0.4230769230769231, 0.875, 'sc <= -0.318\ngini = 0.469\nsamples = 280\nvalue = [175, 105]\nnclass = notckd'),
Text(0.15384615384615385, 0.625, 'sg <= 1.017\ngini = 0.351\nsamples = 132\nvalue = [30, 102]\nnclass = notckd'),
Text(0.07692307692307693, 0.375, 'gini = 0.0\nsamples = 23\nvalue = [23, 0]\nnclass = notckd'),
Text(0.23076923076923078, 0.375, 'al <= 0.5\ngini = 0.12\nsamples = 109\nvalue = [7, 102]\nnclass = notckd'),
Text(0.15384615384615385, 0.125, 'gini = 0.019\nsamples = 103\nvalue = [1.0, 102.0]\nnclass = notckd'),
Text(0.3076923076923077, 0.125, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]\nnclass = notckd'),
Text(0.6923076923076923, 0.625, 'sg <= 1.022\ngini = 0.04\nsamples = 148\nvalue = [145, 3]\nnclass = notckd'),
Text(0.5384615384615384, 0.375, 'age <= 1.605\ngini = 0.014\nsamples = 140\nvalue = [139, 1]\nnclass = notckd'),
Text(0.46153846153846156, 0.125, 'gini = 0.0\nsamples = 136\nvalue = [136, 0]\nnclass = notckd'),
Text(0.6153846153846154, 0.125, 'gini = 0.375\nsamples = 4\nvalue = [3, 1]\nnclass = notckd'),
Text(0.8461538461538461, 0.375, 'age <= -0.407\ngini = 0.375\nsamples = 8\nvalue = [6, 2]\nnclass = notckd'),
Text(0.7692307692307693, 0.125, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]\nnclass = ckd'),
Text(0.9230769230769231, 0.125, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]\nnclass = notckd')]

```



```

pred_new = cs_dt_new.predict(X_test)
pred_new[:5]

```

```
array([0, 0, 0, 0, 0], dtype=int8)
```



```
cm_new = pd.DataFrame(confusion_matrix(y_test, pred_new), index=['No', 'Yes'], columns=['No',
cm_new.index.name = 'True'
cm_new.columns.name = 'Predicted'
cm_new
```

	Predicted	
	No	Yes
True		
No	72	3
Yes	1	44

```
print(classification_report(y_test, pred_new))
```

	precision	recall	f1-score	support
0	0.99	0.96	0.97	75
1	0.94	0.98	0.96	45
accuracy			0.97	120
macro avg	0.96	0.97	0.96	120
weighted avg	0.97	0.97	0.97	120

Algorithm 2: KNN


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import neighbors
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn import metrics
```

```
x2_train, x2_test, y2_train, y2_test = train_test_split(
    X1_filled,
    y,
    test_size = 0.3,
    random_state = 0,
    stratify = y
)
```

```
k_range = range(1, 15)
scores_r = []

for k in k_range:
    knn_r = neighbors.KNeighborsClassifier(n_neighbors=k)
    y_train_flat = np.ravel(y2_train)
    knn_r.fit(x2_train, y_train_flat)
    y_r_pred = knn_r.predict(x2_test.values)
    scores_r.append(round(metrics.accuracy_score(y2_test, y_r_pred), 2))
```

```
plt.plot(k_range, scores_r)
plt.xlabel('K')
plt.ylabel('test set accuracy')
plt.xticks(range(1,15))
plt.show()
```



```
knn3 = neighbors.KNeighborsClassifier(  
    n_neighbors = 2,  
    algorithm = 'brute'  
)
```

```
y_train_flat = np.ravel(y2_train)
```

```
knn3.fit(x2_train, y_train_flat)
```

```
KNeighborsClassifier(algorithm='brute', n_neighbors=2)
```

```
pred3 = knn3.predict(x2_test.values)
```

```
pred3[1:5]
```

```
array([0, 0, 0, 0], dtype=int8)
```

```
round(metrics.accuracy_score(y2_test, pred3), 2)
```

0.99

```
conf_matrix = confusion_matrix(y2_test, pred3)
print(conf_matrix)
```

```
[[74  1]
 [ 0 45]]
```

10. **Performance Metrics:** Outline the two metrics for comparing the performance of the classifiers.

```
cm = pd.DataFrame(confusion_matrix(y_test, pred), index=['No', 'Yes'], columns=['No', 'Yes'])
cm.index.name = 'True'
cm.columns.name = 'Predicted'
print('Confusion Matrix of Decision Tree:')
cm
```

Confusion Matrix of Decision Tree:

	Predicted	
	No	Yes
True	No	Yes
	72	3
Yes	1	44

For the decision tree classifier confusion matrix, there are 72 observations been correctly predicted as No, and 44 variables been predicted as Yes truly. Reminded that the y variable has two input, one is no, means this person do not have chronic kidney disease; the other one is yes, means the person has chronic kidney disease.

```

conf_matrix = pd.DataFrame(confusion_matrix(y2_test, pred3), index=['No', 'Yes'], columns=['No', 'Yes'])
conf_matrix.index.name = 'True'
conf_matrix.columns.name = 'Predicted'
print('Confusion Matrix of KNN:')
conf_matrix

```

Confusion Matrix of KNN:

	Predicted	
True	No	Yes
	No	Yes
No	74	1
Yes	0	45

73 observations without chronic kidney disease in KNN classifier been coreectly predicted as no, and 45 person with chronic kidney disease are correctly predicted as yes.

By comparision of these two confussion matrixs, since the KNN classifier has number of observations of both person with chronic kidney disease and person without chronic kidney disease been predicted truely than the number of observations been predicted by the decision tree classifier. Thus, we considered KNN classifier better than decision tree. Moreover, compare the accuracy score of both classifiers, KNN has 98% accuracy which is greater than decision tree with 97% accuracy.

11. **Feature Selection/Extraction:** Implement methods to enhance the performance of at least one classifier in (9). The answer for this question can be included in (12).

We want to do feature selection for decision tree.


```

feature_importances = cs_dt.feature_importances_

sorted_indices = feature_importances.argsort()[::-1]
sorted_feature_names = X_train.columns[sorted_indices]
sorted_importances = feature_importances[sorted_indices]

```

```
sns.barplot(x = sorted_importances, y = sorted_feature_names)
plt.show()
```



12. **Classifier Comparison:** Utilize the selected metrics to compare the classifiers based on the test set. Discuss your findings (at least two statements).

```
conf_matrix = confusion_matrix(y2_test, pred3)
print("KNN Result: ")
print(conf_matrix)

print("Decision Tree Result: " )
cm = pd.DataFrame(confusion_matrix(y_test, pred), index=['No', 'Yes'], columns=['No', 'Yes'])
cm.index.name = 'True'
cm.columns.name = 'Predicted'
print(cm)
```

KNN Result:

```
[[74  1]
 [ 0 45]]
```

Decision Tree Result:

Predicted	No	Yes
True		
No	72	3
Yes	1	44

From the above confusion matrix, we can observe that KNN has a better result.

13. **Interpretable Classifier Insight:** After re-training the interpretable classifier with all available data, analyze and interpret the significance of predictor variables in the context of the data and the challenge (at least two statements).

```
clf = DecisionTreeClassifier(max_depth =6,
    random_state=1,criterion='entropy')
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
```

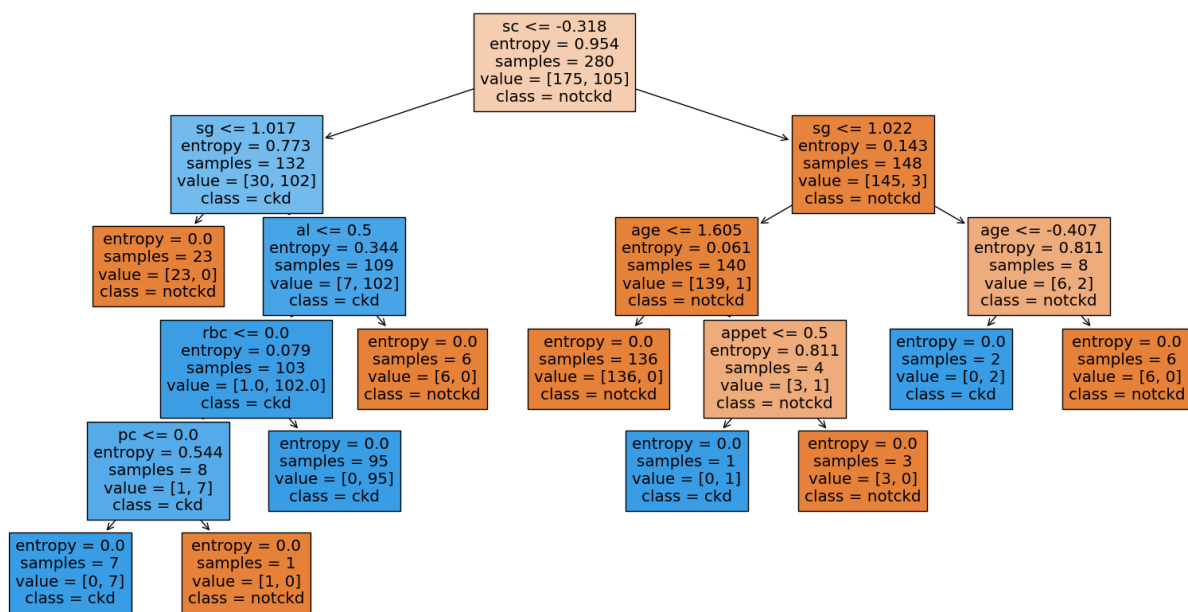
```
from sklearn.metrics import accuracy_score, classification_report

print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))

plt.figure(figsize=(20,10))
plot_tree(clf, filled=True,feature_names = X_train.columns.tolist(),class_names=['notckd', 'ckd'])
plt.show()
```

Accuracy: 0.9666666666666667

	precision	recall	f1-score	support
0	0.99	0.96	0.97	75
1	0.94	0.98	0.96	45
accuracy			0.97	120
macro avg	0.96	0.97	0.96	120
weighted avg	0.97	0.97	0.97	120



```


cm = pd.DataFrame(confusion_matrix(y_test, y_pred), index=['No', 'Yes'], columns=['No', 'Yes'])
cm.index.name = 'True'
cm.columns.name = 'Predicted'
cm

```

	Predicted No	Predicted Yes
True No	72	3
True Yes	1	44


```
fea_imp = clf.feature_importances_
sorted_indices = fea_imp.argsort()[::-1]
sorted_feature_names = X_train.columns[sorted_indices]
sorted_importances = fea_imp[sorted_indices]
sns.barplot(x = sorted_importances, y = sorted_feature_names)
plt.ylabel("variables")

plt.show()
```



We have retrain the model with a different criteria and less depth. This less complicated model give as a better result. We also find features that are significant to the model.

14. **[Bonus]** Sub-group Improvement Strategy: If sub-groups were identified, propose and implement a method to improve one classifier performance further. Compare the performance of the new classifier with the results in (12).

We can use resampling to improve one classifier performance further, since we have too many missing values in this sample data.

```
poly = PolynomialFeatures(degree=3, interaction_only=True)
X_polytrain = poly.fit_transform(X_train)
X_polytest = poly.transform(X_test)
logit_model = LogisticRegression()
logit_model.fit(X_polytrain, y_train)
y_polypred = logit_model.predict(X_polytest)
accuracy = accuracy_score(y_test, y_polypred)
accuracy
```

0.9833333333333333

Improved Model Accuracy is 0.9833333333333333.

15. **Team Contributions:** Document each team member's specific contributions related to the questions above.

Ling Dai(400220798): Question Question 1-6,8,9,14

Shujia Zhang(400390028): Question 2, 5, 7,10,11,12,13

Yunxin Li(400323756):3,4, 10, 11, 12, 13

All team memebers are participated in this project actively. We have a lot of disscusion and communication during the projects, so most of the code are actually done together offline in school.

16. **Link** to the public GitHub repository.

<https://github.com/dail9/3da-assignment6/blob/main/assignment6.ipynb>

Reference

Rubini, Soundarapandian, L., and P. Eswaran. 2015. "Chronic Kidney Disease." UCI Machine Learning Repository.